SalesForge AI

Phased Build Roadmap

v1.0 · 2024–2025 · Internal Engineering & Product

Phase 1 8–12 Weeks MVP Launch	Phase 2 16–20 Weeks Intelligence Layer	Phase 3 24+ Weeks Enterprise Scale

This document defines the complete engineering and product roadmap for SalesForge AI — a B2B enterprise sales intelligence platform. It covers sprint-by-sprint deliverables, feature specifications, technical architecture, success KPIs, risk registers, and team structure across all three phases of development.

Executive Summary

SalesForge AI is built in three deliberate phases. Phase 1 establishes a shippable MVP delivering core value — lead enrichment, AI-generated outreach, a CRM pipeline, and a performance dashboard — in 8–12 weeks with a lean team. Phase 2 adds the intelligence layer that differentiates the product: meeting AI, call coaching, deal scoring, and org mapping. Phase 3 scales the platform to enterprise grade with SSO, advanced forecasting, a market intelligence feed, and a public API ecosystem.

Strategic Principles

- Ship fast, validate early — no feature built without a success metric
- Claude API is the AI backbone — every intelligence feature routes through it
- Database-first — schema is the source of truth, UI and APIs derive from it
- Multi-tenant from day one — every table scoped to organization_id
- Ruthless MVP scoping — Phase 1 ships one thing well per module

Three-Phase Summary

Phase	Timeline	Team Size	Core Theme	Gate to Next Phase
1 — MVP				
2 — Intelligence	Wks 13–20	5–6 devs	AI differentiation at depth	NPS > 40, 80% retention
3 — Enterprise				

Phase 1: MVP Launch Weeks 1–12

Lead enrichment · AI outreach · Basic CRM · Pipeline dashboard

Phase 1 focuses on one objective: get a sales rep from zero to sending a personalized, AI-generated outreach email to a qualified lead — all within a single product session. Everything else is cut. The CRM is deliberately simple. The dashboard shows only what matters. Speed of shipping beats depth of features.

Phase 1 Feature Specification

Feature	What It Does	Tech / API	Priority
ICP Builder	Form-based ideal customer profile creator with industry, size, geo, revenue, tech stack filters	PostgreSQL + React form	**P0**
Lead Enrichment	Ingest ICP → query Explorium API → return scored, enriched lead list	Explorium API + pgvector	**P0**
Pain Hypothesis Generator	Claude analyzes company data to generate "why they need you" per lead	Claude API (claude-sonnet-4-6)	**P0**
AI Email Composer	3-touch cold email sequence per lead with tone controls and inline editing	Claude API + React editor	**P0**
LinkedIn DM Generator	2-touch LinkedIn sequence with character limit enforcement	Claude API	**P1**
Email Sending	One-click send via connected Gmail/Outlook, open + click tracking	SendGrid + OAuth	**P0**
Pipeline Kanban	6-stage deal board: Prospecting → Closed. Drag-and-drop stage changes	React DnD + PostgreSQL	**P0**
Deal Cards	Company, contact, value, last activity, days in stage, AI next action badge	React + REST API	**P0**
Activity Log	Manual + auto-logged emails, calls, notes per deal and lead	PostgreSQL activities table	**P0**
Stall Detection	Flag deals inactive 7+ days. Visual indicator on kanban card	DB trigger + cron job	**P1**
Performance Dashboard	Pipeline value by stage, quota tracker, weekly activity counts, top 5 deals	React charts + views	**P0**

Feature	What It Does	Tech / API	Priority
Multi-tenant Auth	Org-scoped login, role-based access (admin/manager/rep/viewer)	Clerk + JWT + RLS	**P0**
Team Management	Invite users, assign leads/deals, basic admin controls	REST API + email invite	**P1**

Phase 1 Sprint Plan

Sprint	Focus	Deliverables	Success Metric
Sprint 1 Wks 1–2			
Sprint 2 Wks 3–4	Lead Engine	ICP builder form · Explorium API integration · lead enrichment pipeline · lead list UI with scoring · pagination + filters	Generate 100 enriched leads from a test ICP in < 30s
Sprint 3 Wks 5–6			
Sprint 4 Wks 7–8	Email Delivery	Gmail/Outlook OAuth integration · SendGrid fallback · open/click tracking pixels · sequence scheduling · bounce handling · unsubscribe link	0 emails land in spam during internal testing
Sprint 5 Wks 9–10			
Sprint 6 Wks 11–12	Dashboard + Polish	Performance dashboard · quota tracker · rep activity metrics · onboarding flow · error states · loading states · beta user testing · bug bash	10 beta reps complete full workflow without hand-holding

Phase 1 Claude Integration Architecture

Every AI call in Phase 1 routes through the Claude API with structured context injection:

Feature	Prompt Strategy	Model	Avg Tokens
Email Sequence	System: Tone guide + proven email frameworks (AIDA, PAS) User: Pain hypothesis + lead role + sequence step number	Sonnet 4.6	~1,500
Next Best Action	System: Deal stage definitions + sales playbook User: Deal stage, days inactive, last activity type, stakeholder map	Sonnet 4.6	~1,000

Phase 1 KPIs

KPI	Target	Measurement Method
Lead enrichment accuracy	> 90% correct company/contact data	Manual audit of 100 random leads
Reply rate (AI-generated)	> 8%	SendGrid + manual tagging
Beta user session length	> 15 minutes avg	Analytics (PostHog/Mixpanel)

Phase 1 Risk Register

Risk	Likelihood	Impact	Mitigation
Explorium API rate limits throttle enrichment	**Medium**	**High**	Implement request queue + exponential backoff. Cache enrichment results for 30 days. Build CSV import fallback.

Risk	Likelihood	Impact	Mitigation
Email deliverability — AI emails flagged as spam	**Medium**	**High**	Warm up sending domain over 2 weeks. Use SendGrid dedicated IP. Human review + editing before send.
Claude API latency > 10s for sequence generation	**Low**	**Medium**	Show streaming output. Cache pain hypothesis for 7 days per lead. Pre-generate for top-scored leads.
Scope creep delays MVP to Week 16+	**High**	**High**	Strict P0/P1/P2 gating. Weekly scope review. PM has kill-switch authority on new feature requests.
Multi-tenant RLS misconfiguration causes data leak	**Low**	**Critical**	PostgreSQL Row Level Security on all tables. Automated cross-tenant penetration test in CI pipeline.

Phase 2: Intelligence Layer Weeks 13–20

Meeting AI · Call coaching · Org mapping · Deal scoring · Competitive intelligence

Phase 2 is where SalesForge AI stops being a better CRM and starts being an unfair advantage. Meeting intelligence captures and synthesizes every customer conversation. Call coaching gives real-time guidance during live calls. Deal scoring predicts outcomes before reps can feel them. Org mapping reveals the hidden power structures that win or lose enterprise deals. This is the moat.

Phase 2 Feature Specification

Feature	What It Does	Tech / API	Priority
Meeting Transcription	Real-time transcription via Whisper API or Recall.ai bot. Store raw transcript linked to deal	Recall.ai / AssemblyAI	**P0**
AI Meeting Summary	Claude generates structured summary: key points, objections raised, commitments made, sentiment	Claude API	**P0**
Action Item Extraction	Auto-extract action items from transcript with owner, description, due date. Push to tasks table	Claude API + NLP	**P0**
CRM Auto-Update from Meeting	After meeting: auto-update deal stage, next steps, last activity. One-click approval flow	Claude API + REST	**P1**
Live Call Coaching	Real-time suggestions during call: when to ask discovery question, objection detected, talk:listen ratio alert	WebSocket + Claude streaming	**P1**
Post-Call Scorecard	Score call on: discovery quality, objection handling, next steps secured, rapport. Coaching tips	Claude API	**P1**
AI Deal Scoring	Win probability 0–100 based on: stage velocity, stakeholder coverage, engagement depth, competitor presence	ML model + Claude explanation	**P0**
Deal Risk Flags	Auto-flag: single-threaded deal, no economic buyer contact, proposal > 14 days no response	Rule engine + Claude analysis	**P0**
Stakeholder Org Mapper	Visual org chart per deal: champion, blocker, economic	React Flow + PostgreSQL	**P0**

Feature	What It Does	Tech / API	Priority
	buyer, influencer. Relationship health per contact		
Multi-threading Assistant	Claude suggests new contacts to reach at target account based on deal stage and missing roles	Claude API + Explorium	**P1**
Competitive Battlecards	Auto-generated battlecards from win/loss data. Trigger when competitor name appears in meeting notes	Claude API + competitors table	**P1**
Email Reply Intelligence	Classify inbound replies: interested / objection / referral / not now. Suggest response	Claude API + email webhook	**P0**
Manager Deal Review AI	Manager view: AI-flagged deals needing attention, rep coaching opportunities, forecast accuracy	Claude API + dashboard	**P1**
Sequence A/B Testing	Test two email variants against each other. Statistical significance calculator. Auto-promote winner	PostgreSQL + analytics	**P2**

Phase 2 Sprint Plan

Sprint	Focus	Deliverables	Success Metric
Sprint 7 Wks 13–14			
Sprint 8 Wks 15–16	Deal Intelligence	AI deal scoring model · win probability display on kanban · risk flag engine (single-threaded, no EB, stale proposal) · deal health dashboard widget	Score accuracy > 70% on held-out historical deals
Sprint 9 Wks 17–18			
Sprint 10 Wks 19–20	Call Coaching + Reply Intel	WebSocket live transcription pipeline · real-time coaching nudge UI · post-call scorecard · email reply classifier	80% of reps use call coaching feature weekly within 30 days

Sprint	Focus	Deliverables	Success Metric
		· response suggester · Phase 2 QA + user testing	

Phase 2 Claude Integration Architecture

Feature	Prompt Strategy	Mode	Latency Target
Action Items	System: Extract only confirmed commitments with owners User: Meeting summary + participant map	Batch	< 15s
Live Call Coaching	System: Real-time sales coach. Short, actionable, one sentence max User: Last 30 seconds of transcript + deal context	Streaming	< 2s

Phase 2 KPIs

KPI	Target	Measurement Method
Deal score predictive accuracy	> 70% correct win/loss prediction at proposal stage	Backtested on 3 months of historical data
Risk flag false positive rate	< 15% of flagged deals are false alarms	Rep feedback + outcome tracking
NPS after Phase 2	> 40	In-app NPS survey at 60 days

Phase 2 Risk Register

Risk	Likelihood	Impact	Mitigation
Live call coaching latency > 3s feels disruptive	**Medium**	**High**	Use Claude streaming API. Pre-warm context at call start. Debounce triggers to 30s intervals. A/B test opt-in vs. always-on.
Meeting bot rejected by enterprise security policies	**Medium**	**Medium**	Offer Zoom/Teams native integration as fallback. Support manual transcript upload. Build browser extension option.
Deal scoring model biased by small training data	**High**	**Medium**	Start with rule-based scoring. Layer ML only after 500+ historical deals. Be transparent about confidence intervals.
Competitive battlecards legally sensitive	**Low**	**Medium**	Legal review of competitive content policy. User-generated claims only. No factual accuracy guarantees in UI.
Phase 2 scope bloats — team burns out	**High**	**High**	Enforce 2-week sprint reviews. CTO has final authority on scope. Hire 1–2 additional engineers before Sprint 7.

Phase 3: Enterprise Scale Weeks 21–36+

Market intelligence · Advanced forecasting · Enterprise SSO · Public API ecosystem · Compliance

Phase 3 transforms SalesForge AI from a product reps love into a platform enterprises buy. This means SSO, SCIM provisioning, SOC 2 compliance, SLA-backed uptime, advanced forecasting that finance teams trust, a market intelligence feed that surfaces opportunities before competitors see them, and a public API that lets enterprise customers build on top of SalesForge infrastructure.

Phase 3 Feature Specification

Feature	What It Does	Tech / API	Priority
Market Intelligence Feed	Daily digest: funding rounds, leadership changes, hiring surges, product launches at target accounts. Relevance-ranked, AI outreach suggestion per trigger	Claude + news APIs + Explorium webhooks	**P0**
Advanced Revenue Forecasting	AI-powered quarterly forecast with confidence intervals. Scenario modeling (best/base/worst). Variance explanations. Finance-grade export	Claude + statistical model + PostgreSQL	**P0**
Enterprise SSO (SAML 2.0)	Okta, Azure AD, Google Workspace SSO. SCIM user provisioning. JIT account creation	WorkOS or Auth0 Enterprise	**P0**
SCIM User Provisioning	Automatic user creation/deactivation synced from IdP. Group-based role assignment	WorkOS SCIM	**P0**
SOC 2 Type II Compliance	Audit logging, data residency options, encryption at rest/transit, vendor assessments, penetration testing	Vanta + AWS compliance controls	**P0**
Advanced Analytics Suite	Cohort analysis, rep ramp curves, territory performance, deal velocity benchmarks, win/loss pattern analysis	Metabase or custom charts	**P1**
Public REST API	Full CRUD API for leads, deals, activities. Webhook subscriptions. API key management. Rate limiting. Docs via Swagger	FastAPI + API Gateway	**P0**
Zapier / Make Integration	Pre-built Zaps: new lead → Slack, deal closed → HubSpot,	Zapier developer platform	**P1**

Feature	What It Does	Tech / API	Priority
	meeting → Notion. Trigger and action support		
Salesforce Bi-directional Sync	Two-way sync: leads, contacts, opportunities, activities. Conflict resolution rules. Field mapping UI	Salesforce REST API + webhooks	**P0**
HubSpot Bi-directional Sync	Same as Salesforce for HubSpot CRM	HubSpot API v3	**P1**
Territory Management	Assign geographic/industry/account territories to reps. Auto-route enriched leads. Territory performance reporting	PostgreSQL + admin UI	**P1**
AI Conversation Library	Searchable library of winning call transcripts and emails. Filter by deal size, industry, stage. Rep onboarding use case	pgvector semantic search + Claude	**P1**
Custom AI Playbooks	Org-specific sales playbooks injected into every Claude prompt. Methodology support: MEDDIC, SPICED, Challenger, BANT	Prompt management system	**P0**
Mobile App (iOS + Android)	Activity logging, deal updates, call coaching, notifications. Offline-capable for call notes	React Native + Expo	**P2**
Admin Governance Console	Usage analytics, API key management, data export, user audit log, GDPR data deletion, seat management	React + admin API	**P0**

Phase 3 Sprint Plan

Sprint	Focus	Deliverables	Success Metric
Sprint 11–12 Wks 21–24			
Sprint 13–14 Wks 25–28	Market Intelligence Feed	News API integrations (Crunchbase, LinkedIn, Google News) · trigger detection pipeline · relevance scoring · Claude outreach suggestion · daily digest email · in-app feed UI	100 relevant triggers surfaced per week per org with > 70% relevance rating

Sprint	Focus	Deliverables	Success Metric
Sprint 15–16 Wks 29–32			
Sprint 17–18 Wks 33–36	Public API + Ecosystem	REST API v1 · Swagger docs · API key management · rate limiting · Zapier integration · webhook subscriptions · developer portal · SDK (Python + JS) · mobile app MVP	3 design partners using API in production. 500 Zapier connects in first month.

Phase 3 KPIs

KPI	Target	Measurement Method
Forecast accuracy at quarter close	Within 15% of actuals	Forecasted vs. actual revenue delta
SOC 2 Type II certification	Achieved by Week 32	Vanta audit completion
Market triggers actioned	> 40% of surfaced triggers result in outreach	Trigger action tracking in DB
Churn rate	< 5% monthly	Subscription cancellation rate

Phase 3 Risk Register

Risk	Likelihood	Impact	Mitigation
SOC 2 audit takes longer than planned — blocks enterprise deals	**Medium**	**High**	Start Vanta setup in Week 21. Hire compliance consultant. Use "SOC 2 in progress" letter for sales. Target Type I first.
Salesforce API complexity delays CRM sync quality	**High**	**Medium**	Partner with a Salesforce SI. Use Fivetran as interim ETL. Scope sync to 5 core objects initially.

Risk	Likelihood	Impact	Mitigation
Public API abuse or DDoS attack	Low	High	Rate limiting at API gateway (100 req/min free, 1000 req/min paid). WAF on AWS. Anomaly detection alerts.
Revenue forecasting accuracy disappoints finance buyers	Medium	High	Set expectations: forecast is a guide, not a guarantee. Show confidence intervals. Release only after 2 quarter backtest.
Competitor launches similar feature set mid-Phase 3	Medium	High	Ship market intel feed early (Sprint 13). Invest in Claude-native features competitors cannot easily replicate.

Full Technology Stack

Layer	Technology	Purpose	Phase Introduced
	Next.js 14 + React 18		
UI Library	Tailwind CSS + shadcn/ui	Component system, design tokens	Phase 1
	Zustand + React Query		
Charts	Recharts + D3.js	Dashboard analytics visualizations	Phase 1
	React DnD Kit		
Org Charts	React Flow	Stakeholder org mapper	Phase 2
	Node.js + Fastify		
ORM	Drizzle ORM	Type-safe DB queries, migrations	Phase 1
	PostgreSQL 15		
Vector DB	pgvector extension	Semantic search for conversation library	Phase 3
	Redis (Upstash)		
Queue	BullMQ on Redis	Async jobs: enrichment, email, AI calls	Phase 1
	Clerk		
Enterprise Auth	WorkOS	SAML SSO, SCIM provisioning	Phase 3
	Claude API (Sonnet 4.6)		
Transcription	Recall.ai + AssemblyAI	Meeting bots + real-time transcription	Phase 2
	SendGrid		

Layer	Technology	Purpose	Phase Introduced
Enrichment	Explorium API	Lead and company data enrichment	Phase 1
	Stripe		
Infra	AWS (ECS Fargate + RDS)	Production hosting, managed Postgres	Phase 1
	CloudFront + S3		
Monitoring	Datadog + Sentry	APM, error tracking, alerting	Phase 1
	PostHog		
Compliance	Vanta	SOC 2 Type II automation	Phase 3
	GitHub Actions		
Public API	AWS API Gateway	Rate limiting, API key auth, routing	Phase 3

Team Structure by Phase

Role	Phase 1 Focus	Phase 2 Focus	Phase 3
Frontend Engineer (1)	Kanban, dashboard, email composer	Org mapper, call coaching UI	Mobile app
AI/Prompt Engineer (1)	Claude prompt design + evaluation	Live coaching, meeting summary	Playbook system
QA / DevOps (0.5)	CI/CD, test coverage	Load testing, monitoring	SOC 2 controls

Phase Gate Criteria

Each phase has binary gate criteria. ALL must be met before engineering resources shift to the next phase. No partial credit.

Gate	Criteria (ALL required)	Owner
Phase 1 → Phase 2		
Phase 2 → Phase 3	NPS > 40 · Monthly retention > 80% · Deal score accuracy > 65% · Call coaching used by > 50% of reps weekly · $50K MRR or 1 enterprise LOI	Board review
Phase 3 Success		